

Lab9

Lily Lau

20/03/2021

```
seeds <- read.table(
  "https://archive.ics.uci.edu/ml/machine-learning-databases/00236/seeds_dataset.txt"
)
colnames(seeds) <- c("area",
  "perimeter",
  "compactness",
  "length_of_kernel",
  "width_of_kernel",
  "asy_coeff",
  "length_of_kernel_groove",
  "Class")
summary(seeds)
```

```
##      area      perimeter      compactness      length_of_kernel
## Min.   :10.59   Min.   :12.41   Min.   :0.8081   Min.   :4.899
## 1st Qu.:12.27   1st Qu.:13.45   1st Qu.:0.8569   1st Qu.:5.262
## Median :14.36   Median :14.32   Median :0.8734   Median :5.524
## Mean   :14.85   Mean   :14.56   Mean   :0.8710   Mean   :5.629
## 3rd Qu.:17.30   3rd Qu.:15.71   3rd Qu.:0.8878   3rd Qu.:5.980
## Max.   :21.18   Max.   :17.25   Max.   :0.9183   Max.   :6.675
## width_of_kernel  asy_coeff      length_of_kernel_groove      Class
## Min.   :2.630   Min.   :0.7651   Min.   :4.519           Min.   :1
## 1st Qu.:2.944   1st Qu.:2.5615   1st Qu.:5.045           1st Qu.:1
## Median :3.237   Median :3.5990   Median :5.223           Median :2
## Mean   :3.259   Mean   :3.7002   Mean   :5.408           Mean   :2
## 3rd Qu.:3.562   3rd Qu.:4.7687   3rd Qu.:5.877           3rd Qu.:3
## Max.   :4.033   Max.   :8.4560   Max.   :6.550           Max.   :3
```

```
cor(dplyr::select(seeds, -Class))
```

```
##              area  perimeter  compactness  length_of_kernel
## area              1.0000000  0.9943409    0.6082884    0.9499854
## perimeter         0.9943409  1.0000000    0.5292436    0.9724223
## compactness       0.6082884  0.5292436    1.0000000    0.3679151
## length_of_kernel  0.9499854  0.9724223    0.3679151    1.0000000
## width_of_kernel   0.9707706  0.9448294    0.7616345    0.8604149
## asy_coeff        -0.2295723 -0.2173404   -0.3314709    -0.1715624
## length_of_kernel_groove 0.8636927  0.8907839    0.2268248    0.9328061
##              width_of_kernel  asy_coeff  length_of_kernel_groove
## area              0.9707706 -0.22957233    0.86369275
```

```
## perimeter          0.9448294 -0.21734037          0.89078390
## compactness        0.7616345 -0.33147087          0.22682482
## length_of_kernel   0.8604149 -0.17156243          0.93280609
## width_of_kernel    1.0000000 -0.25803655          0.74913147
## asy_coeff          -0.2580365  1.00000000         -0.01107902
## length_of_kernel_groove 0.7491315 -0.01107902          1.00000000
```

- Print tidy seeds data set.

```
dim(seeds)
```

```
## [1] 210  8
```

```
knitr::kable(head(seeds)) %>%
  kable_styling(latex_options="scale_down")
```

area	perimeter	compactness	length_of_kernel	width_of_kernel	asy_coeff	length_of_kernel_groove	Class
15.26	14.84	0.8710	5.763	3.312	2.221	5.220	1
14.88	14.57	0.8811	5.554	3.333	1.018	4.956	1
14.29	14.09	0.9050	5.291	3.337	2.699	4.825	1
13.84	13.94	0.8955	5.324	3.379	2.259	4.805	1
16.14	14.99	0.9034	5.658	3.562	1.355	5.175	1
14.38	14.21	0.8951	5.386	3.312	2.462	4.956	1

- We scale the predictors.

```
x <- seeds %>%
  dplyr::select(-Class) %>%
  scale()
```

- We split 75%/25% training/test set.

```
set.seed(1)

seeds_train_index <- seeds %>%
  mutate(ind = 1:nrow(seeds)) %>%
  group_by(Class) %>%
  mutate(n = n()) %>%
  sample_frac(size = .75, weight = n) %>%
  ungroup() %>%
  pull(ind)
```

- We create binary output units $y_k, k = 1, 2, 3$ using `class.ind()`.

```
library(nnet)
class_labels <- pull(seeds, Class) %>%
  class.ind()
knitr::kable(head(class_labels)) %>%
  kable_styling(latex_options="scale_down")
```

[illegible]

- Create predictor matrix for training/test set and output for training/test set.

```
seeds_train <- x[seeds_train_index, ]
train_class <- class_labels[seeds_train_index,]
seeds_test <- x[-seeds_train_index, ]
test_class <- class_labels[-seeds_train_index,]
```

- Let's look at the help page for `nnet()`.
- Let's tune `size` = number of units in the hidden layer and `decay` = weight decay parameter.

```
nn_seeds <- nnet(
  x = seeds_train,
  y = train_class,
  size = 4,
  decay = 0,
  softmax = TRUE,
  maxit=500
)
```

```
## # weights:  47
## initial  value 179.079752
## iter   10 value 10.357187
## iter   20 value  0.304073
## iter   30 value  0.002143
## iter   40 value  0.000138
## iter   40 value  0.000061
## iter   40 value  0.000061
## final   value  0.000061
## converged
```

- Compute test error for NN with `size` = 4 and `decay` = 0.

```
nn_pred <- predict(nn_seeds, seeds_test,
  type="class")

tab_seeds <- table(slice(
  seeds,
  -seeds_train_index) %>% pull(Class),
  nn_pred)

1-sum(diag(tab_seeds))/sum(tab_seeds)
```

```
## [1] 0.1111111
```

Neural networks (Boston data (quantitative response))

- Let's consider housing price data, *Boston* in MASS package.
- Response is quantitative.

```
library(nnet)
library(MASS)
```

```
##
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':
##
##      select
```

- We scale predictors and response.
- We split training/test set.

```
train_Boston <- sample(
  1:nrow(Boston),
  nrow(Boston)/2
)

x <- scale(Boston)
```

- Create predictor matrix for training/test set and output for training/test set.

```
Boston_train <- x[train_Boston, ]
train_medv <- x[train_Boston, "medv"]
Boston_test <- x[-train_Boston, ]
test_medv <- x[-train_Boston, "medv"]
```

- Let's tune `size` = number of units in the hidden layer and `decay` = weight decay parameter.

```
nn_Boston <- nnet(
  Boston_train,
  train_medv,
  size=10,
  decay=1,
  softmax=FALSE,
  maxit=1000,
  linout=TRUE
)
```


```
## # weights: 161
## initial value 469.211580
## iter 10 value 39.116735
## iter 20 value 22.164051
## iter 30 value 17.626264
## iter 40 value 14.619830
## iter 50 value 12.655570
## iter 60 value 11.292161
## iter 70 value 10.583592
## iter 80 value 10.254760
## iter 90 value 10.097962
```

```
## iter 100 value 10.015590
## iter 110 value 9.948224
## iter 120 value 9.917840
## iter 130 value 9.905889
## iter 140 value 9.901823
## iter 150 value 9.900874
## iter 160 value 9.900509
## iter 170 value 9.900410
## iter 180 value 9.900337
## iter 190 value 9.900141
## iter 200 value 9.900086
## iter 210 value 9.900065
## final value 9.900062
## converged
```

Compute test error for the above model: NN with `size = 10` and `decay = 1`.

```
nn_pred <- predict(
  nn_Boston,
  Boston_test,
  type="raw"
)
```

```
plot(test_medv, nn_pred)
```



```
mean((test_medv - nn_pred)^2)
```

```
## [1] 0.003687532
```

CV for NN - Iris data

- 80%/20% training/test set.

Let's vary the size of hidden layer- Iris data

```
library(e1071)
library(cluster)
set.seed(1)

data("iris")

Species <- pull(iris, Species)

xy <- dplyr::select(iris, -Species) %>%
  scale() %>%
  data.frame() %>%
  mutate(Species = Species) # scale predictors

iris_train_index <- iris %>%
  mutate(ind = 1:nrow(iris)) %>%
  group_by(Species) %>%
  mutate(n = n()) %>%
  sample_frac(size = .8, weight = n) %>%
  ungroup() %>%
  pull(ind)

iris_train <- slice(xy, iris_train_index)
iris_test <- slice(xy, -iris_train_index)
class_labels <- pull(xy, Species) %>%
  class.ind()

iris_nnet1 <- tune.nnet(
  Species~.,
  data = iris_train,
  size = 1:30,
  tunecontrol = tune.control(sampling = "cross", cross=5)
)

head(summary(iris_nnet1))

## $best.parameters
##   size
## 2     2
##
## $best.performance
```


```

## [1] 0.03333333
##
## $method
## [1] "nnet"
##
## $nparcomb
## [1] 30
##
## $train.ind
## $train.ind$(0.881,24.8]
## [1] 40 83 90 35 111 112 120 78 22 70 28 37 61 46 67 71 116 44 49
## [20] 117 56 89 50 7 20 100 80 99 16 2 118 65 79 101 41 77 107 13
## [39] 109 114 82 19 17 57 11 31 115 74 95 55 45 52 68 119 9 97 81
## [58] 113 108 85 32 87 94 12 30 14 62 6 72 64 38 102 91 3 104 69
## [77] 54 5 88 33 84 47 8 4 98 18 27 36 63 110 25 21 66 73 23
## [96] 75
##
## $train.ind$(24.8,48.6]
## [1] 106 96 103 60 51 93 34 10 1 43 59 26 15 58 29 24 42 48 76
## [20] 39 105 53 92 86 20 100 80 99 16 2 118 65 79 101 41 77 107 13
## [39] 109 114 82 19 17 57 11 31 115 74 95 55 45 52 68 119 9 97 81
## [58] 113 108 85 32 87 94 12 30 14 62 6 72 64 38 102 91 3 104 69
## [77] 54 5 88 33 84 47 8 4 98 18 27 36 63 110 25 21 66 73 23
## [96] 75
##
## $train.ind$(48.6,72.4]
## [1] 106 96 103 60 51 93 34 10 1 43 59 26 15 58 29 24 42 48 76
## [20] 39 105 53 92 86 40 83 90 35 111 112 120 78 22 70 28 37 61 46
## [39] 67 71 116 44 49 117 56 89 50 7 95 55 45 52 68 119 9 97 81
## [58] 113 108 85 32 87 94 12 30 14 62 6 72 64 38 102 91 3 104 69
## [77] 54 5 88 33 84 47 8 4 98 18 27 36 63 110 25 21 66 73 23
## [96] 75
##
## $train.ind$(72.4,96.2]
## [1] 106 96 103 60 51 93 34 10 1 43 59 26 15 58 29 24 42 48 76
## [20] 39 105 53 92 86 40 83 90 35 111 112 120 78 22 70 28 37 61 46
## [39] 67 71 116 44 49 117 56 89 50 7 20 100 80 99 16 2 118 65 79
## [58] 101 41 77 107 13 109 114 82 19 17 57 11 31 115 74 91 3 104 69
## [77] 54 5 88 33 84 47 8 4 98 18 27 36 63 110 25 21 66 73 23
## [96] 75
##
## $train.ind$(96.2,120]
## [1] 106 96 103 60 51 93 34 10 1 43 59 26 15 58 29 24 42 48 76
## [20] 39 105 53 92 86 40 83 90 35 111 112 120 78 22 70 28 37 61 46
## [39] 67 71 116 44 49 117 56 89 50 7 20 100 80 99 16 2 118 65 79
## [58] 101 41 77 107 13 109 114 82 19 17 57 11 31 115 74 95 55 45 52
## [77] 68 119 9 97 81 113 108 85 32 87 94 12 30 14 62 6 72 64 38
## [96] 102
##
##
## $sampling
## [1] "5-fold cross validation"

```



```
plot(iris_nnet1)
```



Model selection - size = `iris_nnet1$best.parameters[1,1]` with relatively small CV error and dispersion of CV error.

Fit the model with size = `iris_nnet1$best.parameters[1,1]`.

```
library(nnet)
nn_iris <- nnet(
  x = dplyr::select(iris_train, -Species),
  y = class_labels[iris_train_index, ],
  size = iris_nnet1$best.parameters[1,1],
  decay = 0,
  softmax = TRUE
)
```

```
## # weights:  19
## initial  value 139.787195
## iter  10 value 51.659855
## iter  20 value 12.382653
## iter  30 value  2.538123
## iter  40 value  0.820028
## iter  50 value  0.000596
## iter  60 value  0.000139
## final   value  0.000086
## converged
```

Compute test error for the selected model with size = iris_nnet1\$best.parameters[1,1]

```
nn_pred <- predict(
  nn_iris,
  dplyr::select(iris_test, -Species),
  type="class"
)

tab <- table(pull(iris_test, Species),
  nn_pred
)

tab
```

```
##           nn_pred
##           setosa versicolor virginica
## setosa           10           0           0
## versicolor        0          10           0
## virginica         0           2           8
```

```
1- sum(diag(tab))/sum(tab)
```

```
## [1] 0.06666667
```

Let's tune decay and size - Iris data

```
set.seed(1)

iris_nnet2 <- tune.nnet(
  Species~.,
  data = iris_train,
  size = 1:20,
  decay = 0:3,
  tunecontrol = tune.control(sampling = "cross", cross=5)
)

head(summary(iris_nnet2))
```

```
## $best.parameters
##   size decay
## 11   11    0
##
## $best.performance
## [1] 0.01666667
##
## $method
## [1] "nnet"
##
## $nparcomb
## [1] 80
```

```

##
## $train.ind
## $train.ind$(0.881,24.8] '
## [1] 99 44 102 33 84 35 70 105 42 38 20 28 86 95 90 40 83 25 113
## [20] 119 111 88 6 24 32 114 2 45 18 22 65 13 81 94 48 63 23 46
## [39] 92 77 29 66 67 56 101 80 62 93 69 108 31 116 17 9 57 60 19
## [58] 26 30 72 53 110 10 118 11 27 75 15 50 103 91 16 47 12 104 112
## [77] 8 49 3 98 64 55 71 96 36 4 115 5 52 41 61 120 78 58 107
## [96] 76
##
## $train.ind$(24.8,48.6] '
## [1] 68 39 1 34 87 43 14 82 59 51 97 85 21 106 54 74 7 73 79
## [20] 109 37 89 100 117 32 114 2 45 18 22 65 13 81 94 48 63 23 46
## [39] 92 77 29 66 67 56 101 80 62 93 69 108 31 116 17 9 57 60 19
## [58] 26 30 72 53 110 10 118 11 27 75 15 50 103 91 16 47 12 104 112
## [77] 8 49 3 98 64 55 71 96 36 4 115 5 52 41 61 120 78 58 107
## [96] 76
##
## $train.ind$(48.6,72.4] '
## [1] 68 39 1 34 87 43 14 82 59 51 97 85 21 106 54 74 7 73 79
## [20] 109 37 89 100 117 99 44 102 33 84 35 70 105 42 38 20 28 86 95
## [39] 90 40 83 25 113 119 111 88 6 24 69 108 31 116 17 9 57 60 19
## [58] 26 30 72 53 110 10 118 11 27 75 15 50 103 91 16 47 12 104 112
## [77] 8 49 3 98 64 55 71 96 36 4 115 5 52 41 61 120 78 58 107
## [96] 76
##
## $train.ind$(72.4,96.2] '
## [1] 68 39 1 34 87 43 14 82 59 51 97 85 21 106 54 74 7 73 79
## [20] 109 37 89 100 117 99 44 102 33 84 35 70 105 42 38 20 28 86 95
## [39] 90 40 83 25 113 119 111 88 6 24 32 114 2 45 18 22 65 13 81
## [58] 94 48 63 23 46 92 77 29 66 67 56 101 80 62 93 47 12 104 112
## [77] 8 49 3 98 64 55 71 96 36 4 115 5 52 41 61 120 78 58 107
## [96] 76
##
## $train.ind$(96.2,120] '
## [1] 68 39 1 34 87 43 14 82 59 51 97 85 21 106 54 74 7 73 79
## [20] 109 37 89 100 117 99 44 102 33 84 35 70 105 42 38 20 28 86 95
## [39] 90 40 83 25 113 119 111 88 6 24 32 114 2 45 18 22 65 13 81
## [58] 94 48 63 23 46 92 77 29 66 67 56 101 80 62 93 69 108 31 116
## [77] 17 9 57 60 19 26 30 72 53 110 10 118 11 27 75 15 50 103 91
## [96] 16
##
##
## $sampling
## [1] "5-fold cross validation"


```

```

plot(iris_nnet2)

```

Performance of `nnet`



We can choose `size = iris_nnet2$best.parameters[1,1]` and `decay = iris_nnet2$best.parameters[1,2]`.

Fit the model with `size = iris_nnet2$best.parameters[1,1]` and `decay = iris_nnet2$best.parameters[1,2]`.

```
nn_iris_d_s <- nnet(  
  x = dplyr::select(iris_train, -Species),  
  y = class_labels[iris_train_index, ],  
  size = iris_nnet2$best.parameters[1,1],  
  decay = iris_nnet2$best.parameters[1,2],  
  softmax = TRUE  
)
```

```
## # weights:  91  
## initial  value 164.446139  
## iter  10 value 15.814895  
## iter  20 value 1.891497  
## iter  30 value 0.102615  
## final   value 0.000056  
## converged
```

```
# Compute test error  
nn_pred <- predict(  
  nn_iris_d_s,  
  dplyr::select(iris_test, -Species),  
  type="class"  
)
```

```
tab <- table(pull(iris_test, Species),
  nn_pred
)
```

```
tab
```

```
##           nn_pred
##           setosa versicolor virginica
##   setosa         10          0         0
##   versicolor      0          10         0
##   virginica       0           2         8
```

```
1- sum(diag(tab))/sum(tab)
```

```
## [1] 0.06666667
```

Clustering -coffee data

K-means clustering assignment depends on the initial cluster assignments. Thus, we need to run the clustering with different random assignment and select the best solution (the clustering with the minimum total within sum of squares).

Coffee - from the help page - data on the chemical composition of coffee samples collected from around the world, comprising 43 samples from 29 countries. We dropped the first two columns of the data.

k-means -coffee data

- Let's apply k-means for clustering.

```
library(cluster)
library(factoextra) # PCA
```

```
## Warning: package 'factoextra' was built under R version 4.0.4
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(pgmm) # coffee data
data("coffee")
set.seed(1)
x <- dplyr::select(coffee, - Variety, - Country)
x_scaled <- scale(x)
kmeans_coffee <- kmeans(x_scaled, 2)
kmeans_coffee$tot.withinss
```

```
## [1] 330.8912
```

```
kmeans_coffee <- kmeans(x_scaled, 3)
kmeans_coffee$tot.withinss
```

```
## [1] 267.2453
```

```


# Let's select K using elbow method
withclusterss <- function(K,x){
  kmeans(x, K)$tot.withinss
}

K <- 1:8

wcss <- lapply(as.list(K), function(k){
  withclusterss(k, x_scaled)
}) %>% unlist()

ggplot(tibble(K = K, wcss = wcss), aes(x = K, y = wcss)) +
  geom_point() +
  geom_line() +
  xlab("Number of clusters (k)") +
  ylab("Total within-clusters sum of squares") +
  scale_x_continuous(breaks=c(seq(1,K[length(K)])))

```




Based on the elbow method, we can use $k = 2$. There are `dim(x_scaled)[2]` variables. So we need to use dimensionality reduction technique and then plot the clusters in two-dimensions.

```

kmeans_coffee <- kmeans(x_scaled, 2)
fvPCA <- fviz_cluster(kmeans_coffee,
  x_scaled,
  ellipse.type = "norm",

```

```
fvPCA
      main = "Plot the results of k-means clustering after PCA")
```



- Let's use silhouette plots to choose the best number of clusters.
 - Silhouette measures how similar observations are within clusters.
 - Large average silhouette width indicates an appropriate number of clusters.


```
si <- silhouette(kmeans_coffee$cluster, dist(x_scaled))
head(si)
```

```
##      cluster neighbor sil_width
## [1,]      2         1 0.5252373
## [2,]      2         1 0.4346060
## [3,]      2         1 0.4143200
## [4,]      2         1 0.4932787
## [5,]      2         1 0.4632535
## [6,]      2         1 0.4832208
```

```
#average Silhouette width
mean(si[, 3])
```

```
## [1] 0.4186062
```

```
plot(si, nmax= 80, cex.names=0.6, main = "")
```




```
# Let's select K using average Silhouette width
avgSilhouette <- function(K,x) {
  km_cl <- kmeans(x, K)
  sil <- silhouette(km_cl$cluster, dist(x))
  return(mean(sil[, 3]))
}

K <- 2:8

avgSil <- numeric()
for(i in K){
  avgSil[(i-1)] <- avgSilhouette(i, x_scaled)
}

ggplot(tibble(K = K, avgSil = avgSil), aes(x = K, y = avgSil)) +
  geom_point() +
  geom_line() +
  xlab("Number of clusters (k)") +
  ylab("Average silhouette width") +
  scale_x_continuous(breaks=c(seq(1,K[length(K)])))
```

Based on the average Silhouette width, we can use $k = 2$.

k-medoid clustering


- Let's apply k-medoid clustering - coffee data

```
kmedoid_coffee <- pam(x_scaled, 2)
kmedoid_coffee$silinfo$avg.width
```

```
## [1] 0.4186062
```

```
avgSil <- lapply(as.list(2:8), function(k){
  kmedoid_coffee <- pam(x_scaled, k)
  kmedoid_coffee$silinfo$avg.width
}) %>% unlist()

ggplot(tibble(K = 2:8, avgSil = avgSil), aes(x = K, y = avgSil)) +
  geom_point() +
  geom_line() +
  xlab("Number of clusters (k)") +
  ylab("Average silhouette width for k-medoid") +
  scale_x_continuous(breaks=c(seq(1,K[length(K)])))
```



For k-medoid for clustering, based on the average Silhouette width, we can use $K = 2$.

- We can also use Gap statistic to choose k - Reference **Modern Statistics for Modern Biologists** Chapter 5.7.1.

Clustering - votes data

- We will use *votes.repub* in the cluster package.
- Look at the help page for *votes.repub*

k-means

```
data(votes.repub) # from cluster package
votes.repub_scaled <- scale(votes.repub)
votes.repub_kmeans <- kmeans(votes.repub_scaled, 2)
```

Why `kmean()` doesn't work?


Hierarchical clustering - divisive clustering

- Apply divisive clustering


- To divide the selected cluster, the algorithm first looks for its most disparate observation (i.e., which has the largest average dissimilarity to the other observations of the selected cluster)
- This observation initiates the “splinter group”.
- In subsequent steps, the algorithm reassigns observations that are closer to the “splinter group”

```
library(cluster)
library(factoextra)
divisive_votes <- diana(
  votes.repub,
  metric = "euclidean",
  stand = TRUE
)
plot(divisive_votes)
```

Banner of `diana(x = votes.repub, metric = "euclidean", stand`



Dendrogram of `diana(x = votes.repub, metric = "euclidean", stand = "T`



votes.repub
Divisive Coefficient = 0.86

```
cut_divisive_votes <- cutree(as.hclust(divisive_votes), k = 2)
table(cut_divisive_votes) # 8 and 42 group members
```

```
## cut_divisive_votes
## 1 2
## 8 42
```

```
rownames(votes.repub)[cut_divisive_votes == 1]
```


```
## [1] "Alabama"      "Arkansas"      "Florida"       "Georgia"
## [5] "Louisiana"    "Mississippi"   "South Carolina" "Texas"
```

```
# rownames(votes.repub)[cut_divisive_votes == 2]
```

```
#make a nice dendrogram
```

```
fviz_dend(
  divisive_votes,
  cex = 0.5,
  k = 2, # Cut in 2 groups
  palette = "jco", # Color palette
  main = "Dendrogram for votes data (divisive clustering)")
```


Dendrogram for votes data (divisive clustering)



Hierarchical clustering - agglomerative clustering

```
x <- votes.repub %>%
  scale()
hc_vote <- hclust(dist(x), "ward.D")
plot(hc_vote)
```

Cluster Dendrogram



```
#make a nice dendrogram
fviz_dend(
  hc_vote,
  k = 2, # Cut in 2 groups
  cex = 0.5,
  color_labels_by_k = TRUE,
  rect = TRUE,
  main = "Dendrogram for votes data (agglomerative clustering)"
)
```

Dendrogram for votes data (agglomerative clustering)

